



## Data Transfers between User and Kernel Memory via Dedicated DMA Engine

**Candidate:**

Antonio Papa

**Supervisor:**

Prof. Marco Cesati

**Ass. Supervisor:**

Eng. Emiliano Betti

**2013-2014**

# Presentation Outline

- 1 Introduction and Motivation
- 2 DMA-CFTU Extension
- 3 Experimental Results
- 4 Conclusions and Future Works

# Memory Copy Operations

- Memory copy operations are implemented through a sequence of load and store instructions realised by CPU .
- The **load** and **store** instructions use a set of resources as:
  - CPU registers (32 or 64 bit)
  - Cache memory

# Memory Copy Operations

- Memory copy operations are implemented through a sequence of load and store instructions realised by CPU .
- The **load** and **store** instructions use a set of resources as:
  - CPU registers (32 or 64 bit)
  - Cache memory

Furthermore it may happen to have to deal with the **Cache Pollution** problem

## Cache Pollution

An executing computer program loads data into CPU cache unnecessarily, thus causing other useful data to be evicted from the cache into lower levels of the memory hierarchy, degrading performance.

# Memory-Memory operations between kernel and user space

- Every system call makes copy operations between kernel and user space

# Memory-Memory operations between kernel and user space

- Every system call makes copy operations between kernel and user space
- Device drivers have to copy data from/to buffer in user space via *file operations*

# Memory-Memory operations between kernel and user space


- Every system call makes copy operations between kernel and user space
- Device drivers have to copy data from/to buffer in user space via *file operations*
- Furthermore data transfers between Kernel and User space are common in some Kernel sub-system

## **For example:**

- The TCP/IP stack has to copy the payload (*receive-side*) from kernel memory to user application memory

# Data transfers between kernel and user space

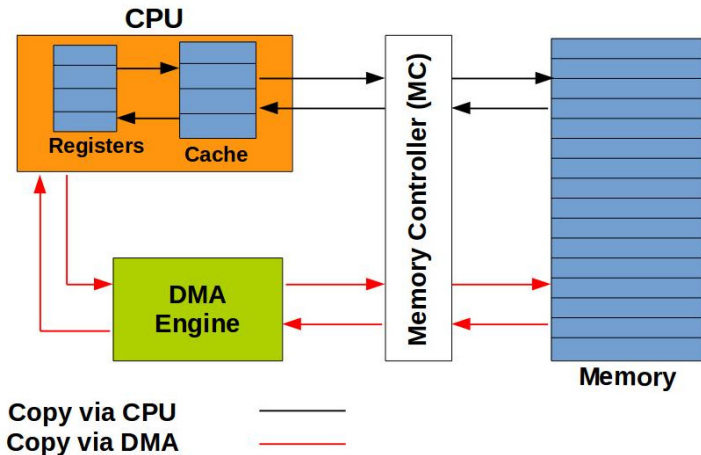
- O.S. exposed to a medium workload
- Time slice of 100 seconds



~ 3,6 M copy operations between kernel and user space



# Motivations for Hardware Copy Engine (1)



# Motivations for Hardware Copy Engine (2)

## **Advantages:**

- Optimisation of CPU resources
- Offload copy operations between kernel and user space

# Motivations for Hardware Copy Engine (2)

## Advantages:

- Optimisation of CPU resources
- Offload copy operations between kernel and user space

## Goals:

- Better performance in data transfer
- Predictability of copy operations
- Control of cache pollution

## Copy memory support

- Generic interface to make DMA transfers
- Kernel copy functions set-up
- Two Policies to manage the DMA channels:
  - *Exclusive Channel*: DMA channel is bound to one CPU (private channel)
  - *Shared Channel*: Every DMA channel is shared (public channel)

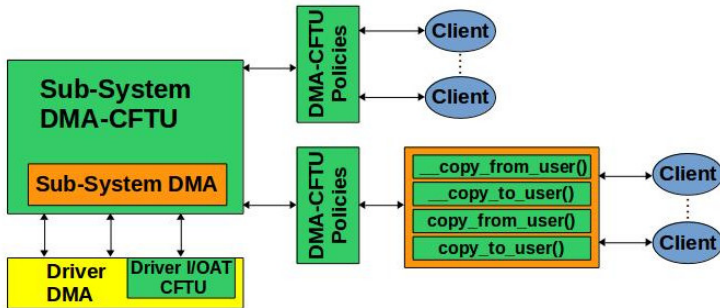
## Copy memory support

- Generic interface to make DMA transfers
- Kernel copy functions set-up
- Two Policies to manage the DMA channels:
  - *Exclusive Channel*: DMA channel is bound to one CPU (private channel)
  - *Shared Channel*: Every DMA channel is shared (public channel)

## Real-time applications support

- Interface to create Real-time policies
- *Priority Channel* Policy


## DMA-CFTU Extension(2)




Components:

- Driver I/OAT-CFTU
- Sub-system DMA-CFTU
- Policies DMA-CFTU

# copy\_to\_user() via DMA engine




# copy\_to\_user() via DMA engine



1 Kernel application calls the `copy_to_user()` function




# copy\_to\_user() via DMA engine




- 1 Kernel application calls the `copy_to_user()` function
- 2 The copy operation is submitted to DMA-CFU System

# copy\_to\_user() via DMA engine



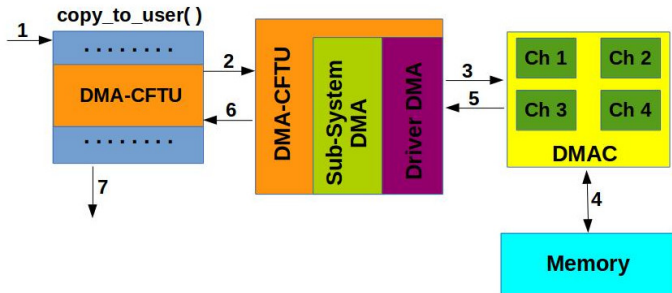
- 1 Kernel application calls the `copy_to_user()` function
- 2 The copy operation is submitted to DMA-CFU System
- 3 Set-up of DMA engine

# copy\_to\_user() via DMA engine




- 1 Kernel application calls the `copy_to_user()` function
- 2 The copy operation is submitted to DMA-CFU System
- 3 Set-up of DMA engine
- 4 The copy operation is performed via DMA engine

# copy\_to\_user() via DMA engine



- 1 Kernel application calls the `copy_to_user()` function
- 2 The copy operation is submitted to DMA-CFU System
- 3 Set-up of DMA engine
- 4 The copy operation is performed via DMA engine
- 5 The DMA engine terminates and sends an interrupt signal

## copy\_to\_user() via DMA engine



- 1 Kernel application calls the `copy_to_user()` function
- 2 The copy operation is submitted to DMA-CFU System
- 3 Set-up of DMA engine
- 4 The copy operation is performed via DMA engine
- 5 The DMA engine terminates and sends an interrupt signal
- 6-7 Return to `copy_to_user()` and termination

# Priority Channel


- *CH-LOW* channels are suspended  $\mathbf{e}_s$  every  $\mathbf{p}_s$
- *CH-HIGH* channels use the DMA bus in a exclusive way  $\mathbf{e}_s$  every  $\mathbf{p}_s$

# Priority Channel

- *CH-LOW* channels are suspended  $e_s$  every  $p_s$
- *CH-HIGH* channels use the DMA bus in a exclusive way  $e_s$  every  $p_s$

For example:

CH-HIGH = ch 1      CH-LOW = ch 2-3-4       $p_s = 90 \text{ us}$        $e_s = 30 \text{ us}$





## Intel's I/O Acceleration Technology (I/OAT)

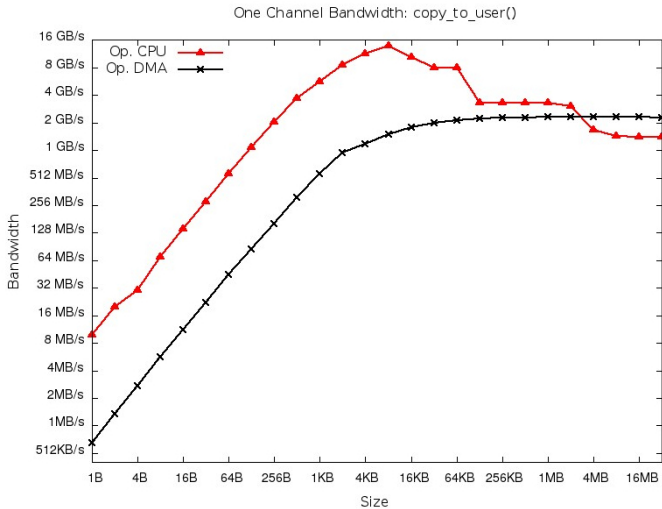
- Asynchronous DMA Copy Engine all'interno dell' MCH
- Memory-Memory Data Transfer
- 4 DMA channels

## System specification

- Intel Xeon 5150 Dual-Core 2.66 GHZ (X2)
- Cache  $L1_{d/i}$  32KB, L2 4MB
- Memory RAM 4GB
- Intel 5000X Chipset MCH with DMA Engine
- Linux Slackware, Kernel version 3.9.2 (patch CFTU)

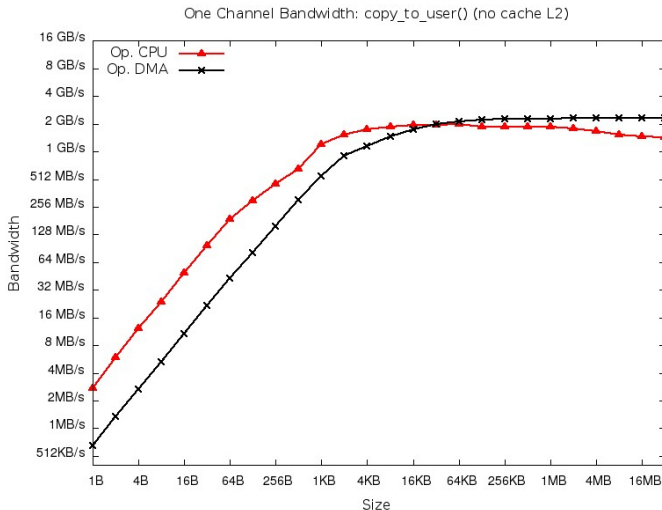


# One Channel Bandwidth (1)



- Buffer > 2MB, the DMA engine has better performance than CPU

# One Channel Bandwidth: (2)




- Buffer > 16 KB, the DMA engine has better performance than CPU

# Priority Channel Bandwidth

**Politica Priority Channel:**  $p_s = 90\mu s$ ,  $e_s = 30\mu s$


- *CH-HIGH* = **chan0** e *CH-LOW* = chan1 , **chan2**, **chan3**



- Buffer > 64 KB: **chan0**: ~ 1.2 GB/s, chan1-2-3: ~ 0.65 GB/s

# Cache Pollution: 16064 Byte Data transfers

- To study the Cache Pollution problem, it has been used the tests of MiBench Benchmark



## Conclusions

- DMA engine performance are better than CPU if the data are not in cache memory
- The DMA transfer rate is not depend on cache memory
- It is possible to control the Cache pollution using the DMA transfer
- It is possible to define scheduling policies to manage the different priorities of DMA channels

## Future Works

- New real-time policies
- Memcopy, Memset And XOR operations via DMA engine
- Processor DMA

# Thank you for your attention!

